

Computation and Communication Co-scheduling for Timely Multi-Task Inference at the Wireless Edge

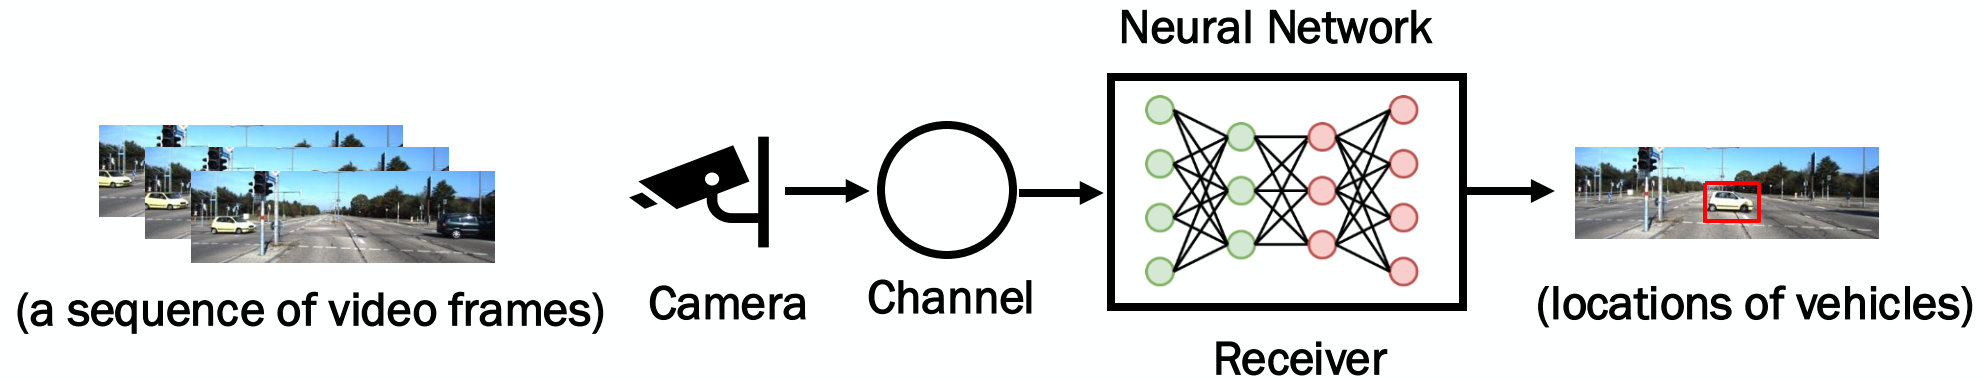
Md Kamran Chowdhury Shisher, Adam Piaseczny, Yin Sun* and
Christopher G. Brinton

Dept. of ECE, Purdue University, USA

Dept. of ECE, Auburn University, USA*




Remote Inference




- A sensor or a camera sends **observed signals** (e.g., a sequence of video frames) to a receiver that predicts a **time-varying target** (e.g., location of vehicles) by using a neural network.
- Due to **limited communication resources**, sending high dimensional sensor observations to a remote receiver is not efficient. The observations can get delayed, and information can be lost.
- This will yield **inaccurate inference** that can affect real-time decision for critical applications.

Remote Inference



- We *split* the neural network to **feature generator** and **predictor**.
- Use feature generator in the edge device which takes the signal observation and generates low dimensional features, then sends them to the receiver.

Multi-Task Inference



- Predict locations of vehicles
- Classify road signs




- Classify friendly Vs. hostile agents
- Predict position of own agents



- Classify customer reaction
- Predict current inventory

- From **autonomous vehicle**, **military**, **smart retail** to **Digital Twin**, edge device may need to perform **multiple inference tasks**.
- Edge device may have **limited computation resources** to **generate features** for all tasks at the same time slot.

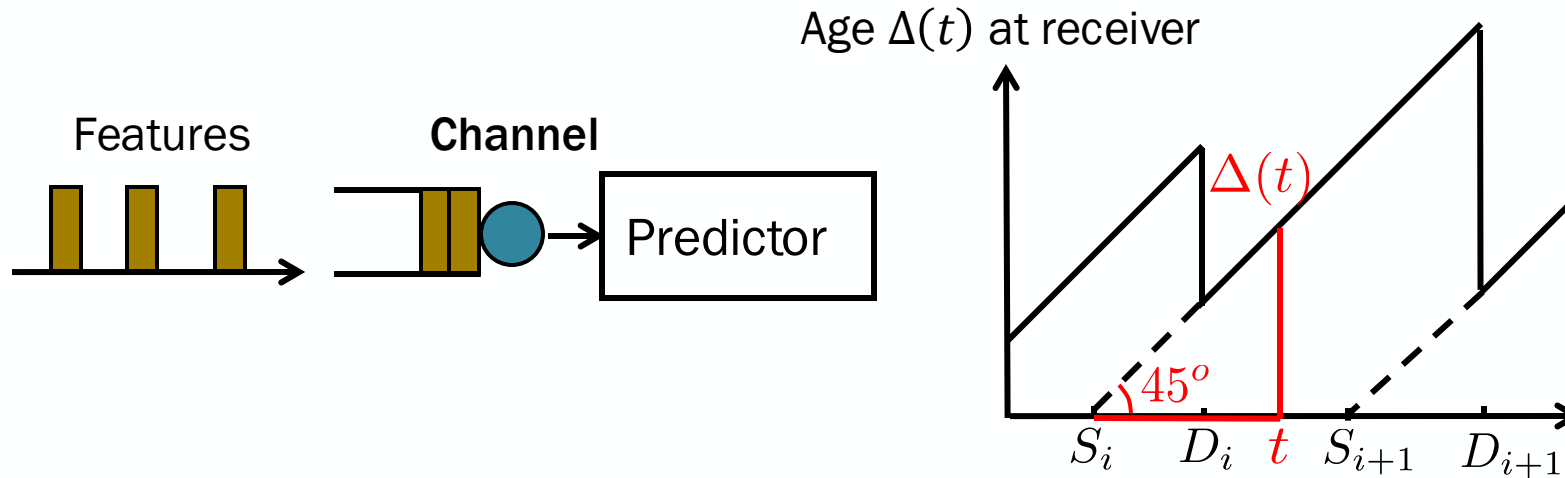
Multi-Task Remote Inference with Multiple Sources



- M edge devices are connected to a receiver
- Receiver predicts K_m targets for each device m based on the most recently delivered features
- Delivered features **may not be fresh** due to **limited computation and communication resources**.
- We use **Age of Information (AoI)** to measure freshness

How can we develop a **computation and communication co-scheduling methodology** to minimize the inference errors across tasks while adhering to network resource constraints?

Age of Information



Definition: At time t , the **Age of Information (AoI)** $\Delta(t)$ is time difference between the current time t and the generation time $t - \Delta(t)$ of the freshest received feature

If feature i is generated at S_i and delivered at D_i

$$\Delta(t) = t - \max\{S_i : D_i \leq t\}$$

Time difference between data generation and usage

Inference Error

Definition: At time t , the **Inference error** for j -th inference task of source m can be expressed as **a function of Aol** $\Delta_{m,j}(t) = \delta$ [Shisher and Sun, MobiHoc' 22, Shisher et al, ToN' 24]

$$p_{m,j}(\delta) = \mathbb{E}_{Y, X \sim P_{Y_{m,j,t}, X_{m,t-\delta}}} \left[L_{m,j}(Y, \psi_{m,j}(\phi_{m,j}(X), \delta)) \right]$$

- Notations:
 - $X_{m,t-\delta}$ denote the signal value generated δ time slots ago from m -th edge devices
 - $\phi_{m,j}(\cdot)$ and $\psi_{m,j}(\cdot)$ are feature generator and predictor functions, respectively
 - $L(y, \hat{y})$ measures the incurred loss when the actual target is y and the predicted value is \hat{y}

Our results can be applied to any loss function $L(y, \hat{y})$. Some examples are: 0-1 loss, quadratic loss, and log loss

Problem Formulation

- Our goal is to **minimize infinite horizon discounted sum of inference errors** for all inference tasks subject to computation and communication resource constraints:

$$\bar{p}_{opt} = \inf_{\pi \in \Pi} \sum_{t=0}^{\infty} \frac{\gamma^t}{K} \sum_{m=1}^M \sum_{j=1}^{k_m} \mathbb{E}_{\pi} [p_{m,j}(\Delta_{m,j}(t))],$$

s.t. $\underbrace{\sum_{j=1}^{k_m} \pi_{m,j}(t) \leq C_m, t = 0, 1, \dots, m = 1, \dots, M,}_{\text{Computation Resource Constraints}}$

$\underbrace{\sum_{m=1}^M \sum_{j=1}^{k_m} \pi_{m,j}(t) \leq N, t = 0, 1, 2, \dots,}_{\text{Communication Resource Constraint}}$

Combinatorial Decision Problem

- $M + 1$ constraints
- M computation constraints, 1 for each device m
- 1 communication constraint shared by all devices
- Weakly Coupled MDP (PSPACE-Hard)**

- $\pi_{m,j}(t) = 1$: feature for j -th inference task of m -th device is generated and transmitted
- At most, N features can be sent at one time slot
- Edge device m can generate features for at most C_m tasks

Related Works

- Aol-based Scheduling:
 - Prior works [Kadota et al, ToN' 18, Shisher and Sun, MobiHoc' 22, Shisher et al, ToN' 24, Tripathi and Modiano ToN'24, Ornee and Sun MobiHoc' 23] considers only one communication constraint
 - Prior works are modeled as RMAB, a special case of weakly coupled MDP
 - Whittle Index Policy are used in RMAB provided that the problem is indexable
- Systematic introduction of the Remote Inference problem [Shisher and Sun, MobiHoc' 22, Shisher et al, ToN' 24]
- Learning and Communications Co-design for Remote Inference [Shisher et al, JSAIT' 23]
- Interpretation of Information Aging on Remote Inference
 - Information-theoretic interpretation of information aging for Markov signals [Sun and Cyr, SPAWC' 18, JCN' 19]
 - Information-theoretic interpretation of information aging for general non-Markov signals [Shisher et al, INFOCOM Aol Workshop'21, Shisher and Sun, MobiHoc' 22, Shisher et al, ToN' 24]
 - AR-model-based analysis/interpretation of information aging [Shisher and Sun, INFOCOM ASol workshop' 24]
 - Experimental results of remote inference [Shisher and Sun, MobiHoc' 22, Shisher et al, ToN' 24, JSAIT' 23]

Q. How to design **Computation and Communication Co-scheduling**?

Lagrangian Primal Dual

- **Primal Problem: (Reoptimized) At every time τ** given AoI value $\Delta_{m,j}(\tau)$, we truncate the problem to T time slots and apply Lagrange multipliers to constraints

$$\begin{aligned}\bar{p}(\boldsymbol{\lambda}, \boldsymbol{\mu}; \tau : T) = & \inf_{\pi \in \Pi} \sum_{t=\tau}^T \sum_{m=1}^M \sum_{j=1}^{k_m} \frac{\gamma^t \mathbb{E}_{\pi} [p_{m,j}(\Delta_{m,j}(t))]}{K} \\ & + \sum_{t=\tau}^T \sum_{m=1}^M \lambda_{m,t} \frac{\gamma^t}{K} \left(\left(\sum_{j=1}^{k_m} \pi_{m,j}(t) \right) - C_m \right) \\ & + \sum_{t=\tau}^T \mu_t \frac{\gamma^t}{K} \left(\left(\sum_{m=1}^M \sum_{j=1}^{k_m} \pi_{m,j}(t) n_{m,j} \right) - N \right),\end{aligned}$$

- **Dual Problem:** We obtain optimal Lagrange Multipliers after solving the dual problem

$$\max_{(\boldsymbol{\lambda}, \boldsymbol{\mu}) \geq 0} \bar{p}(\boldsymbol{\lambda}, \boldsymbol{\mu}; \tau : T)$$

Solution of Primal and Dual

- Solution of Primal Problem:**

Decompose the Primal problem into **per-inference task problem**:

$$\bar{p}_{m,j}(\lambda_m, \mu; \tau : T) = \inf_{\pi_{m,j} \in \Pi_{m,j}} \sum_{t=\tau}^T \gamma^t \mathbb{E}_{\pi_{m,j}} \left[p_{m,j}(\Delta_{m,j}(t)) + \lambda_{m,t} \pi_{m,j}(t) + \mu_t \pi_{m,j}(t) n_{m,j} \right]$$

We solve the problem by dynamic programming:

$$\min_{\pi_{m,j}(t) \in \{0,1\}} \underbrace{Q_{m,j,t}^{\lambda_m, \mu}(\Delta_{m,j}(t), \pi_{m,j}(t))}_{\text{Action Value Function}}$$

- Solution of Dual Problem:**

$$\begin{aligned} \max_{\lambda \geq 0, \mu \geq 0} & \sum_{t=\tau}^T \sum_{m=1}^M \sum_{j=1}^{k_m} V_{m,j,t}^{\lambda_m, \mu}(\Delta_{m,j}(t)) \\ & - \sum_{t=\tau}^T \sum_{m=1}^M \gamma^{t-\tau} \lambda_{m,t} C_m + \sum_{t=\tau}^T \gamma^{t-\tau} \mu_t N \end{aligned}$$

Value Function

Maximum Gain First Policy (Reoptimized)

- **Gain Index:**

$$\alpha_{m,j,t}(\delta) = Q_{m,j,t}^{\lambda_m^*, \mu^*}(\delta, 0) - Q_{m,j,t}^{\lambda_m^*, \mu^*}(\delta, 1)$$

- **At time t , maximize** sum of Gain Indices of all inference tasks subject to the resource constraints

At time t ,

- We iterate through all tasks, ordered by their maximum gain indices.
- If constraints satisfies, we schedule the task

Algorithm 1: Maximum Gain First (MGF) Policy

```
1 for  $t = 0, 1, \dots$  do
2   Update  $\Delta_{m,j}(t)$  for all  $(m, j)$ 
3   Initialize  $\pi_{m,j}(t) \leftarrow 0$  for all  $(m, j)$ 
4   Get  $\lambda^*$  and  $\mu^*$  that maximizes  $\bar{p}(\lambda, \mu; t : T)$ 
5    $\alpha_{m,j} \leftarrow \alpha_{m,j,t}(\Delta_{m,j}(t))$  for all  $(m, j)$ 
6    $C_{m,\text{curr}} \leftarrow 0$  and  $N_{\text{curr}} \leftarrow 0$ 
7    $A(t) \leftarrow \{(m, j) : \alpha_{m,j} > 0\}$ 
8   while  $A(t)$  is not empty do
9      $(m^*, j^*) \leftarrow \arg \max_{m,j} \alpha_{m,j}$ 
10     $c \leftarrow C_{m^*,\text{curr}} + 1$  and  $n \leftarrow N_{\text{curr}} + n_{m^*,j^*}$ 
11    if  $c \leq C_{m^*}$  and  $n \leq N$  then
12      Update  $\pi_{m^*,j^*}(t) \leftarrow 1$ 
13      Update  $C_{m^*,\text{curr}} \leftarrow c$  and  $N_{\text{curr}} \leftarrow n$ 
14     $A(t) = A(t) \setminus (m^*, j^*)$ 
```

Maximum Gain First Policy (Reoptimized)

Theorem: If all Aol functions $p_{m,j}(\delta)$ are bounded and the following holds

$$T \geq \log_{\frac{1}{\gamma}} \left(\sum_{m=1}^M \sqrt{k_m} \right),$$

then the MGF policy is asymptotically optimal as the number of inference tasks per source increases, i.e.,

$$\bar{p}_{\text{MGF}} - \bar{p}_{\text{opt}} = O\left(\frac{1}{\sum_{m=1}^M \sqrt{k_m}}\right)$$

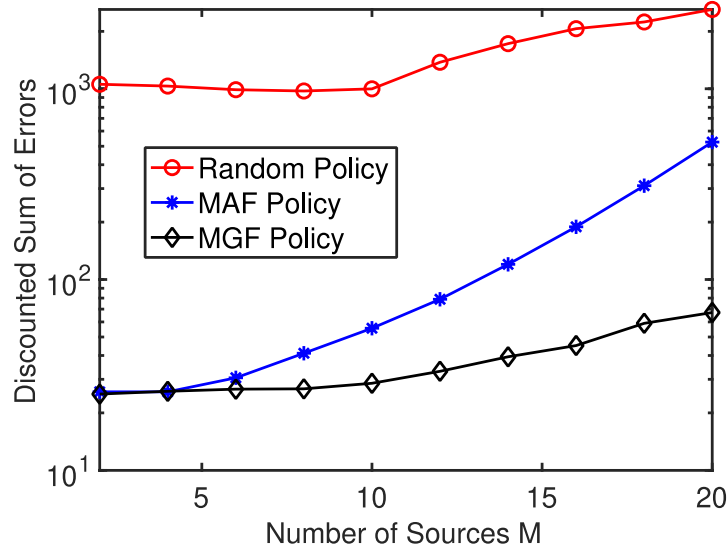
- [\[Brown and Zhang, Operation Research' 23\]](#) introduced Reoptimized Fluid Policy using LP and Occupancy measures. The paper considers all constraints are shared by all sources
- The optimality gap provided in our paper is tighter than

$$\bar{p}_{\text{MGF}} - \bar{p}_{\text{opt}} = O\left(\frac{1}{\sqrt{\sum_{m=1}^M k_m}}\right)$$


of Reoptimized Fluid Policy provided in [\[Brown and Zhang, Operation Research' 23\]](#)

- This is because we use the structural information that M computation constraints are not shared by all devices M

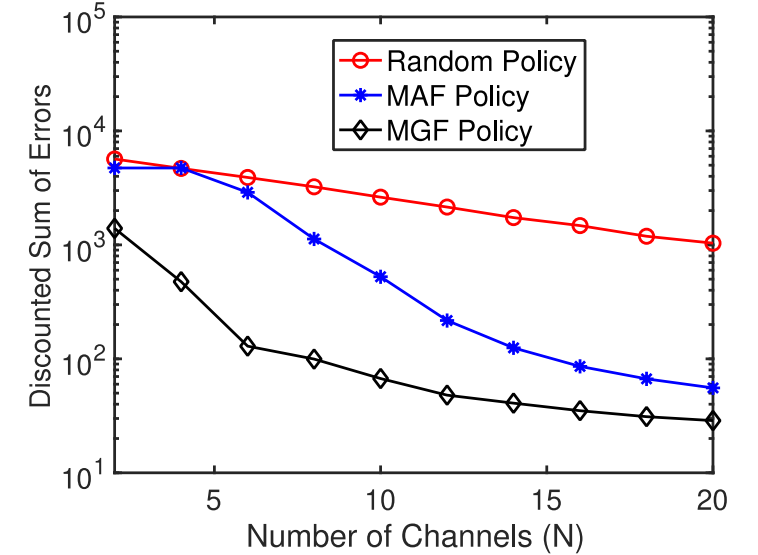
Simulation Results (Synthetic Evaluations)



(a)



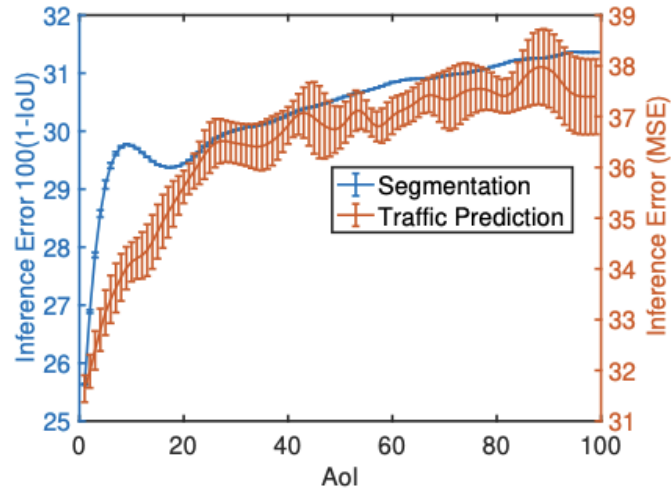
(b)




(c)

- We consider [linear Aol function](#), [logarithmic Aol function](#), and [exponential Aol function](#) [Tripathi and Modiano ToN'24]
- (a) $N = 10$, $k_m = 3$, (b) $M = 20$, $N = 10$, (c) $M = 20$, $k_m = 3$
- Our policy 26 times better compared to Maximum Aol First (**MAF**) and 32 times better compared to **Random Policy**


Simulation Results (Real World Evaluations)



(a) Inference Error vs AoI



(b) Dis. Sum of Errors vs. N



(c) Dis. Sum of Errors vs. M

- We consider **traffic prediction** and **segmentation** using dataset collected from **Next-generation Simulation Program** of **US Department of Transportation Federal Highway Administration**

Summary

Use Inference Error as a metric to design Computation and Communication Co-scheduling

Inference Error = $f(Aol)$

Modeled as Weakly Coupled MDP

Proposed Lagrangian-based Reoptimized (MGF) Policy to solve the problem

- Computation + Communication Policy for optimizing Aol and Remote Inference
 - Remote Inference problem
 - Aol-based Scheduling

Future Works:

- How to optimally design feature generators and predictors?
- Jointly optimize learning, computation, and communication?
- System implementation

Thank You

Email: mshisher@purdue.edu

